

# SSE/AVX SIMD Benchmarking

---

Simone Rossi

May 19, 2017

## 1 INTRODUCTION

In this laboratory, we studied the speedup that Single Instructions Multiple Data (SIMD) instructions (mostly Intel SSE4 and AVX2) can achieve when dealing with vectorial operations.

## 2 USE CASE

Simple use case for this laboratory is a vectorial multiplication:

```
1     for i in 0, ..., N - 1:
2         Z[i] = X[i] * Y[i]
```

The benchmarking has been done by comparing the standard scalar implementation with 128 bits and 256 bits parallelism.

### 3 IMPLEMENTATION

To be able to correctly evaluate the performances in all these three cases, we unrolled the loop exactly eight times per iteration. Let's start by looking the **scalar implementation**.

Listing 1: Scalar Implementation with 8-way loop unrolling

```
1 void componentwise_multiply_real_scalar(int16_t *x, int16_t *y,
    int16_t *z, uint32_t N) {
2     int i;
3     for(i = 0; i < N; i+=8) {
4         z[i] = x[i] * y[i];
5         z[i+1] = x[i+1] * y[i+1];
6         z[i+2] = x[i+2] * y[i+2];
7         z[i+3] = x[i+3] * y[i+3];
8         z[i+4] = x[i+4] * y[i+4];
9         z[i+5] = x[i+5] * y[i+5];
10        z[i+6] = x[i+6] * y[i+6];
11        z[i+7] = x[i+7] * y[i+7];
12    }
13 }
```

Let's see now the use of the SSE intrinsics for the 128 bits parallelism. The right intrinsic for handling the SIMD version of the operation we intend to do is the `__mm_mullo_epi16`, which has the following signature:

```
1 __m128i _mm_mullo_epi16 (__m128i a, __m128i b)
```

According to the Intel Intrinsics Guide, this instruction multiplies the packed 16-bit integers in a and b, producing intermediate 32-bit integers, and store the low 16 bits of the intermediate integers. This is the code to use it.

Listing 2: SSE Implementation with 8-way loop unrolling

```

1  #if defined(__SSE3__) || defined(__SSE4__)
2  void componentwise_multiply_real_sse4(int16_t *x, int16_t *y,
    int16_t *z, uint32_t N) {
3      __m128i *x128 = (__m128i *)x;
4      __m128i *y128 = (__m128i *)y;
5      __m128i *z128 = (__m128i *)z;
6
7      int i = 0;
8      for(i = 0; i < (N>>3); i+=8){
9          z128[i] = _mm_mullo_epi16(x128[i], y128[i]);
10         z128[i+1] = _mm_mullo_epi16(x128[i+1], y128[i+1]);
11         z128[i+2] = _mm_mullo_epi16(x128[i+2], y128[i+2]);
12         z128[i+3] = _mm_mullo_epi16(x128[i+3], y128[i+3]);
13         z128[i+4] = _mm_mullo_epi16(x128[i+4], y128[i+4]);
14         z128[i+5] = _mm_mullo_epi16(x128[i+5], y128[i+5]);
15         z128[i+6] = _mm_mullo_epi16(x128[i+6], y128[i+6]);
16         z128[i+7] = _mm_mullo_epi16(x128[i+7], y128[i+7]);
17     }
18     z = (int16_t*)z128;
19 }
20 #endif

```

Finally, the last one uses AVX instructions: it is basically similar to the previous one, except for the fact the data are packed in 256 bits. The intrinsic to use is `__mm256_mullo_epi16` and the implementation has been adapted as follows:

Listing 3: AVX2 Implementation with 8-way loop unrolling

```

1  #if defined(__AVX2__) && defined(__AVX__)
2  void componentwise_multiply_real_avx2(int16_t *x, int16_t *y,
    int16_t *z, uint32_t N) {
3      __m256i *x256 = (__m256i *)x;

```

```

4     __m256i *y256 = (__m256i *)y;
5     __m256i *z256 = (__m256i *)z;
6
7     int i = 0;
8     for(i = 0; i < (N>>4); i+=8){
9         z256[i] = _mm256_mullo_epi16(x256[i], y256[i]);
10        z256[i+1] = _mm256_mullo_epi16(x256[i+1], y256[i+1]);
11        z256[i+2] = _mm256_mullo_epi16(x256[i+2], y256[i+2]);
12        z256[i+3] = _mm256_mullo_epi16(x256[i+3], y256[i+3]);
13        z256[i+4] = _mm256_mullo_epi16(x256[i+4], y256[i+4]);
14        z256[i+5] = _mm256_mullo_epi16(x256[i+5], y256[i+5]);
15        z256[i+6] = _mm256_mullo_epi16(x256[i+6], y256[i+6]);
16        z256[i+7] = _mm256_mullo_epi16(x256[i+7], y256[i+7]);
17    }
18    z = (int16_t*)z256;
19 }
20 #endif

```

## 4 COMPILATION

For the reason I will mention in the next section, I prefer not to rely on the CLANG compiler but instead I've downloaded and installed the Intel Parallel Studio XE, therefore all the sources have been compiled using the Intel C/C++ Compiler (Version 17.0.4 20170411). Let's have a look on how these routines have been compiled using default -O2 level of optimization.

Listing 4: AVX2 Implementation with 8-way loop unrolling

```

1  _componentwise_multiply_real_sse4:
2  movq    %rdi, %r8
3  shrl    $3, %ecx
4  xorl    %edi, %edi

```

```

5  movq    %rsi, %r9
6  xorl    %esi, %esi
7  testl   %ecx, %ecx
8  jbe     165 <_componentwise_multiply_real_sse4+0xBA>
9  leaq    (%rsi,%r8), %r10
10 vmovdqu (%r10), %xmm0
11 leaq    (%rsi,%rdx), %rax
12 leaq    (%rsi,%r9), %r11
13 vpmullw (%r11), %xmm0, %xmm1
14 addq    $8, %rdi
15 vmovdqu %xmm1, (%rax)
16 vmovdqu 16(%r10), %xmm2
17 vpmullw 16(%r11), %xmm2, %xmm3
18 vmovdqu %xmm3, 16(%rax)
19 vmovdqu 32(%r10), %xmm4
20 vpmullw 32(%r11), %xmm4, %xmm5
21 vmovdqu %xmm5, 32(%rax)
22 vmovdqu 48(%r10), %xmm6
23 vpmullw 48(%r11), %xmm6, %xmm7
24 vmovdqu %xmm7, 48(%rax)
25 vmovdqu 64(%r10), %xmm8
26 vpmullw 64(%r11), %xmm8, %xmm9
27 vmovdqu %xmm9, 64(%rax)
28 vmovdqu 80(%r10), %xmm10
29 vpmullw 80(%r11), %xmm10, %xmm11
30 vmovdqu %xmm11, 80(%rax)
31 vmovdqu 96(%r10), %xmm12
32 vpmullw 96(%r11), %xmm12, %xmm13
33 vmovdqu %xmm13, 96(%rax)

```

```

34 vmovdqu 112(%r10), %xmm14
35 vpmullw 112(%r11), %xmm14, %xmm15
36 addq    $128, %rsi
37 vmovdqu %xmm15, 112(%rax)
38 cmpq    %rcx, %rdi
39 jb      -165 <_componentwise_multiply_real_sse4+0x15>$
40 retq

```

We can see that intrinsics have been replaced with two instructions: `vmovdqu` and `vpmullw`.

The first one (`vmovdqu`) moves 128 bits of packed integer values from the source operand to the destination operand and it is used to load an XMM register from a 128-bit memory location and to store the contents of an XMM register into a 128-bit memory location.

The second one (`vpmullw`), on the other hand, performs a SIMD signed multiply of the packed signed word integers between the two operands, and stores the low 16 bits of each intermediate 32-bit result in the destination operand.

Similar assembly can be found also for the other two functions but they have been omitted to avoid filling the report with pages of assembly (can be found on my GitHub repository of the course).

## 5 PERFORMANCE EVALUATION

All the runs have been done on a laptop with a 4th Generation Intel Core i5 Processor (4308U) running at 2.8 GHz and 16 GB of RAM. Therefore, both SSE4 and AVX2 were available for testing.

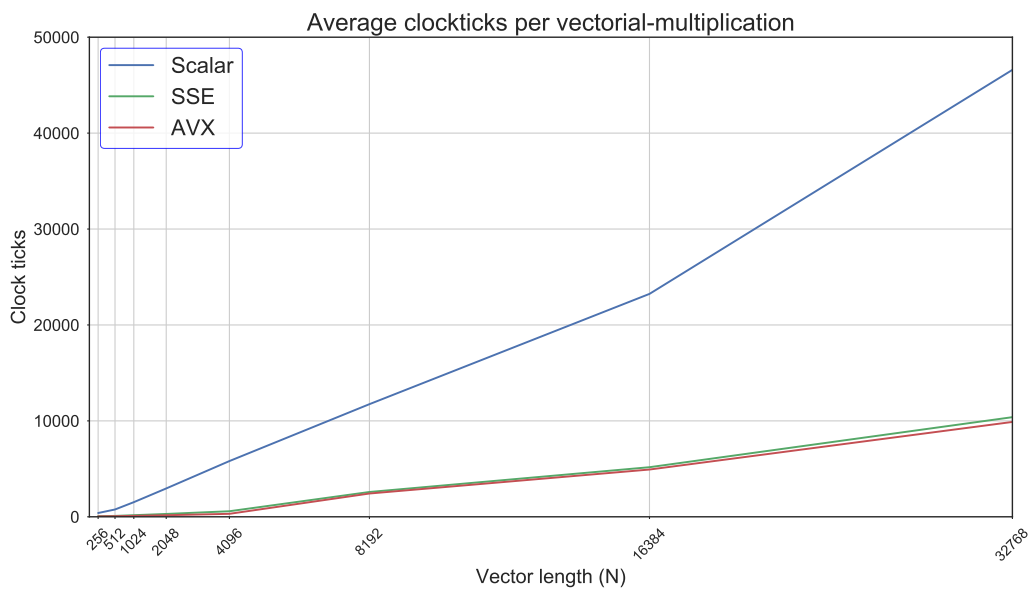
For profiling the performances of these three implementations I used the time measure utilities provided. Unfortunately though I encountered some problems: the binary com-

piled under macOS using CLANG with -O2/3 seemed not to behave correctly in measuring clock ticks. I did not investigate the reason, but instead I decided to modify the way the clock ticks register is read (changing from the instruction `rdtsc` to `__rdtscp`) and using another compiler (Intel C/C++ Compiler).

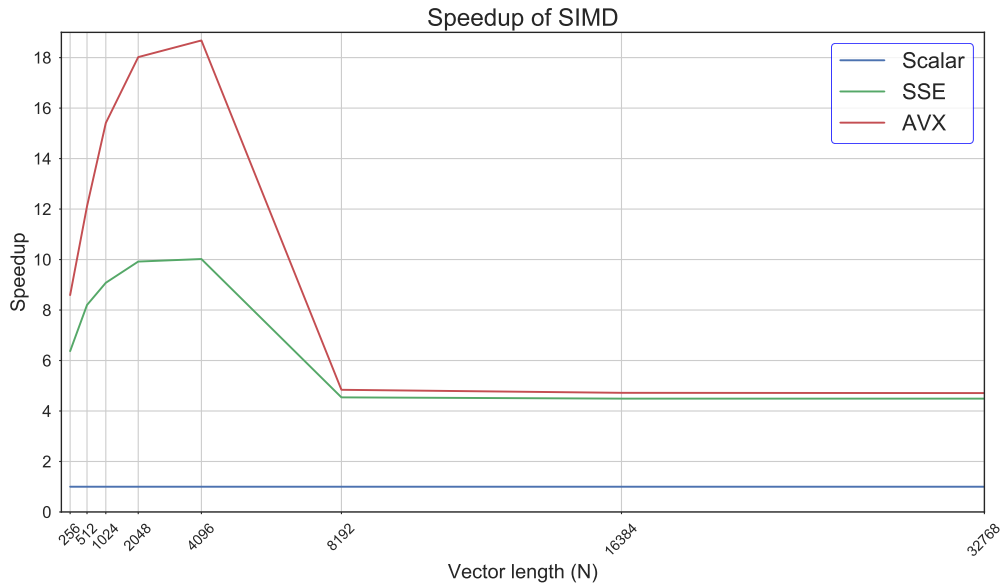
To avoid large bias of cache misses (in the first iterations), each implementation has been tested 10 millions times with a vector lenght spanning from 256 to 32768.

## 6 RESULTS

Let's start by looking at the average ticks per function call, using the default -O2 optimization level.



As we can see, as expected the average clock ticks needed to complete a vectorial multiplication increases linearly with the number of points and we can immediately appreciate the speedup that the SIMD introduces. To better evaluate the performance, let plot the speedup with respect to the scalar baseline.



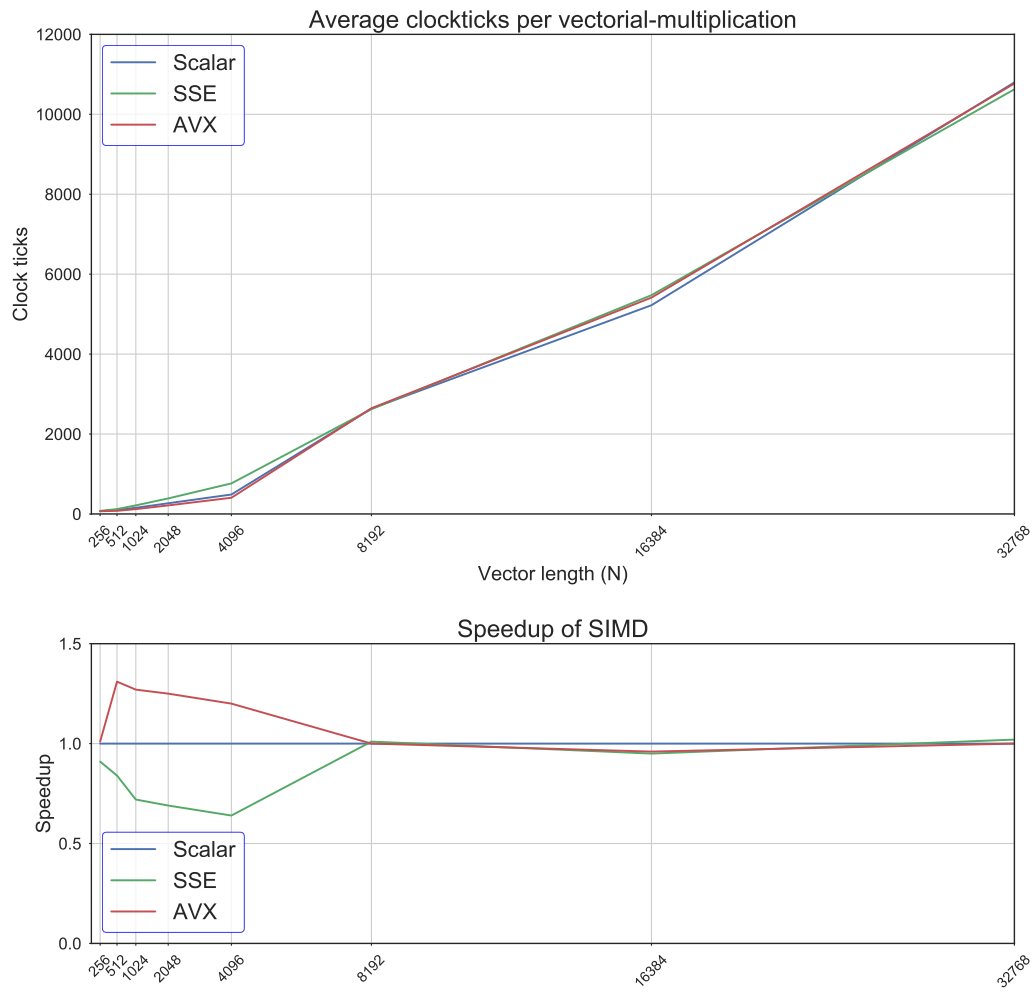
Let's take for instance the case of N equal to 2048. The SSE introduces a 128 bits parallelism and therefore eight 16-bits results can be computed at the same time. We should have expected a plain 8X speedup, but the -O2 optimization was able to reduce the latency by a factor of 10. Similar for AVX, we expected a 16X speedup with respect to the scalar baseline but instead the compiler produced a binary upto 19 times faster than before.

Performances, though, seems to drastically reduce for N equal to 8192 and here we might start having some caches problems: the set of operands arrays and the result array ends up with 384KB of allocated memory which exceeds the available 256KB of L2 cache. Unfortunately, I did not find specifications of latency for the L3 cache and therefore I cannot make comparisons on that.

Compiling the binary with the -O3 flag, ended up with some interesting results.

The compiler was smart enough to understand that the operation we intended to do could have been drastically speeded up using SIMD. Moreover, it's interesting to notice that it automatically used the best version of the available SIMD: in fact, forcing it to use SSE re-





duces the speedup. The same aforementioned caches issues can be found also here starting at 8192 entries.